

# A tutorial on generalised additive mixed effects models for bilingualism research

Stefano Coretta and Joseph V. Casillas

University of Edinburgh | Rutgers University

While recent years have seen a shift towards random effects modelling, particularly in areas of linguistics in which nested structure is the norm (e.g., trial repetitions nested within participants), an over-reliance on standard linear modelling prevails, particularly in the cases of dynamic phenomena that may not constitute a linear relationship, e.g., vowel trajectories, pitch contours, acquisition processes, etc. Generalised Additive (Mixed) Models (GAMMs) are now commonly employed in phonetic research (given the naturally dynamic nature of speech data) and this is reflected by the availability of several tutorials which focus on phonetic data. This tutorial aims at making GAMMs accessible to researchers from other fields within linguistics. In particular, this tutorial is written for researchers in bilingualism and multilingualism who wish to be able to start using GAMMs for non-linear data, which is very common in developmental and learning phenomena. While only the basics will be covered here, we hope that researchers will get the necessary foundations to be able to learn GAMMs from existing resources.

**Keywords:** generalised additive mixed models, GAMMs, non-linear data, dynamic analysis

## 1. Motivation

In recent years quantitative analysis of linguistic data has benefited greatly from the now commonplace use of so-called random effects (also known as hierarchical, mixed-effects, multi-level or nested models (Winter, 2013; Winter, 2020)). While this shift represents a clear improvement in statistical analyses, particularly in areas of linguistics in which nested structure is the norm (e.g., trial repetitions nested within participants), an over-reliance on standard linear modelling prevails, particularly in the cases of dynamic phenomena that may not constitute a linear relationship, e.g., vowel trajectories, pitch contours, acquisition processes, etc.

Generalised Additive (Mixed) Models (GAMMs) are now commonly employed in phonetic research (given the naturally dynamic nature of speech data) and this is reflected by the availability of several tutorials which focus on phonetic data (Sóskuthy, 2017, 2021; Tamminga et al., 2016; Wieling, 2018), and biological data (Pedersen et al., 2019). We hope that this tutorial will make GAMMs accessible to researchers from other fields within linguistics as well. In particular, this tutorial has in mind researchers focusing on bi- and multilingualism who wish to start using GAMMs for non-linear data, which is common in developmental and learning phenomena.

## 1.1 Generalised additive (mixed) models

Generalised Additive Models (GAMs) and their mixed-effects version (GAMMs) represent a variation of the generalised linear model that can accommodate modelling non-linear relationships (Hastie & Tibshirani, 1986; Wood, 2006). The main component specific to GAMMs is known as a smoother or smooth term. These are model terms that can fit non-linear effects. An advantage of using smooth terms is that they are constrained by a smoothing penalty parameter. Importantly, this penalty parameter is a standard part of the fitting procedure that does not depend on user input and can help prevent over-fitting the data. This approach is different from, for example, growth curve analyses which are based on user-specified polynomial degrees. Note, however, that under/over-fitting is still possible depending on the number of basis functions used in building the smoothers (for a more detailed explanation of basis functions, see the aforementioned tutorials).

The rest of this tutorial focuses on applying GAMMs in bilingualism research using the statistical programming language R (R Core Team, 2023). While completing this tutorial will allow you to jump-start into fitting basic GAMMs with the mgcv package Wood (2011), we invite interested readers to peruse the tutorials mentioned above to avoid unnecessary repetition of content here. Readers interested in a more detailed conceptual discussion of GAMMs are referred to Sóskuthy (2017). Finally, we opted not to treat statistical significance testing in GAMMs (significance testing is treated in details in the aforementioned tutorials).

## 2. Pre-requisites

In what follows, we detail the packages needed to complete the tutorial, followed by two specific case studies designed to illustrate how GAMMs can be applied

to model common phenomena in second language acquisition and bilingualism research.

This tutorial assumes readers are already familiar with R (R Core Team, 2023) and have at least some experience fitting linear models, including models with random effects. The following packages need to be installed:

- tidyverse
- mgcv
- tidygam

All the tutorial code and data are freely available here: [https://github.com/stefanocoretta/gamm\\_biling](https://github.com/stefanocoretta/gamm_biling).

### 3. Case study 1: U-shaped learning

#### 3.1 U-shaped learning

Research on language acquisition first documented over-regularization errors in the 60's (Cazden, 1968; Ervin & Miller, 1963). Subsequent research found that over-regularization errors tended to manifest in a developmental curve, and is now referred to as the *U-shaped developmental curve* (Botezatu et al., 2024; Williams et al., 2022). In most cases, U-shaped development is considered to be a three-step process, which begins with accurate performance, followed by a period in which performance dips, and then subsequently becomes accurate again (Carlucci & Case, 2013). Interestingly, U-shaped curves are observed in numerous cognitive-developmental and learning contexts, particularly with regard to child development, e.g., understanding temperature, weight conservation, object permanence, and facial recognition (See Carlucci & Case, 2013), which, in turn, has fostered fruitful research and debate in Cognitive Science regarding learning models (e.g., Marcus et al., 1992; Rumelhart & McClelland, 1986, among others).

In first language acquisition research, the most cited example stems from children's use of the past tense (Cazden, 1968; Ervin & Miller, 1963; Marcus et al., 1992). For instance, Cazden (1968) found that children accurately produced irregular past tense verbs during a period of time, and, at a later point in time, began to produce over-generalization errors, e.g., *feet* followed by *\*feets*. Eventually past tense production again surfaces as expected, typically around the age of three, which coincides with the developmental stage in which children acquire regular past tense verbs (See Marcus et al., 1992). In the realm of Second Language Acquisition (SLA), researchers have documented non-linear, U-shaped courses of development in numerous areas for both developmental and instructional

sequences (See Casani, 2020; Geeslin & Guijarro-Fuentes, 2006; Kellerman, 1985; Long, 1990; Pliatsikas & Marinis, 2013; Shirai, 1990; Williams et al., 2022, among many others). Kellerman (1985), for instance, documented Dutch learners of English acquiring verb alternation patterns of English following a U-shaped trajectory. In a similar vein, Lightbown (1983) showed accuracy producing progressive forms (i.e., *-ing*) followed a U-shaped trajectory when acquiring the simple present/present continuous tense/aspect distinction in French speaking learners of English. U-shaped learning may be a by-product of reorganization or restructuring of prior knowledge (Gass & Selinker, 2008; Shirai, 1990), and likely accounts for language instructors' contention that students often regress when learning new linguistic forms. Be that as it may, it is clear that the non-linear quality of this aspect of language learning requires a non-linear approach to modelling the underlying process.

### 3.2 The data

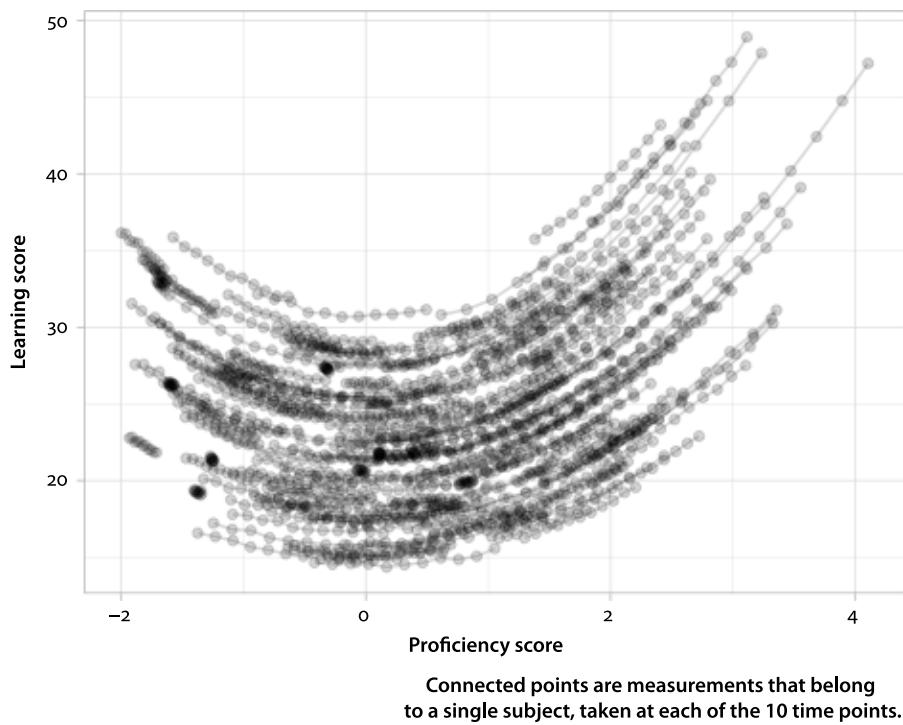
We have simulated data for accuracy scores from 200 language learners, taken at 10 different time points. Note that the simulated data follow a Gaussian (normal) distribution for pedagogical reasons, as it is a straightforward distribution to use when introducing readers to new statistical models. That being said, it is important to recognize that real learner's data 'in the wild' is seldom Gaussian. A proficiency score was also included for each participant at each time point. The data is intended to simulate a study in which participants perform the same learning task and proficiency assessment in a longitudinal design. We print the first six rows of the data frame below. Note that text following a # symbol is a comment and intended to aid the reader in understanding the code.

```
# Load data into R
dat1 <- readRDS("data/dat1.rds")

# Print first 6 rows of the dataframe
head(dat1)

# A tibble: 6 × 4
  score proficiency subj   time
  <dbl>      <dbl> <fct> <int>
1 26.7       -1.37 s1     0
2 26.4       -1.29 s1     1
3 26.0       -1.21 s1     2
4 25.7       -1.14 s1     3
5 25.5       -1.06 s1     4
6 25.2      -0.983 s1     5
```

The score column contains the accuracy scores, while the proficiency column the proficiency scores. The participant ID is given in subj. The time point, o to 9, is in the time column. Figure 1 shows the relationship between proficiency and learning scores for individual participants. From the figure, it is clear that such a relationship has a non-linear, U-shape, in which learning scores initially decrease, plateau, then increase again as proficiency increases.



**Figure 1.** Scatter plot of proficiency and learning scores illustrating a U-shaped learning curve

In the following sections we will analyse this data using GAMMs. For pedagogical purposes, we will first focus on the effect of proficiency scores on learning scores at a single time point, time point 5. Subsequently, we will analyse the entire data set to illustrate how to conduct a time-series analysis.

### 3.3 Modelling a non-linear effect

Let's first focus on how to model a non-linear effect: here, we can look at the effect of the proficiency score on learner accuracy. As we have seen in Figure 1, the effect is U-shaped, i.e., it is not linear. To simplify, we will begin our analysis using data

from a single time point. Later we will expand the analysis to also include time as a predictor. We model a non-linear effect of proficiency on accuracy with the following code.

```
# Subset the data to include only time point 5, assign subset to 'dat15'
dat15 <- filter(dat1, time == 5)

# Load the mgcv package
library(mgcv)

# Fit the model, assign the model object to 'gam_1'
gam_1 <- gam(
  formula = score ~ s(proficiency),
  data = dat15
)
```

The formula states that `score`, the outcome variable (also known as the dependent variable or criterion) is modelled as a function of `proficiency`. Note that we use `s()` to indicate that we want to estimate a (potentially) non-linear effect of proficiency. The name of the function, `s`, stands for “smooth term”. Smooth terms (aka smoothers) are mathematical functions that allow GAMs to fit non-linear effects. A detailed treatment of smoothers is beyond the scope of this tutorial. We refer readers to Wood (2006) for a mathematical description and to other existing tutorials for a gentler introduction (Pedersen et al., 2019; e.g., Sóskuthy, 2017; Wieling, 2018).

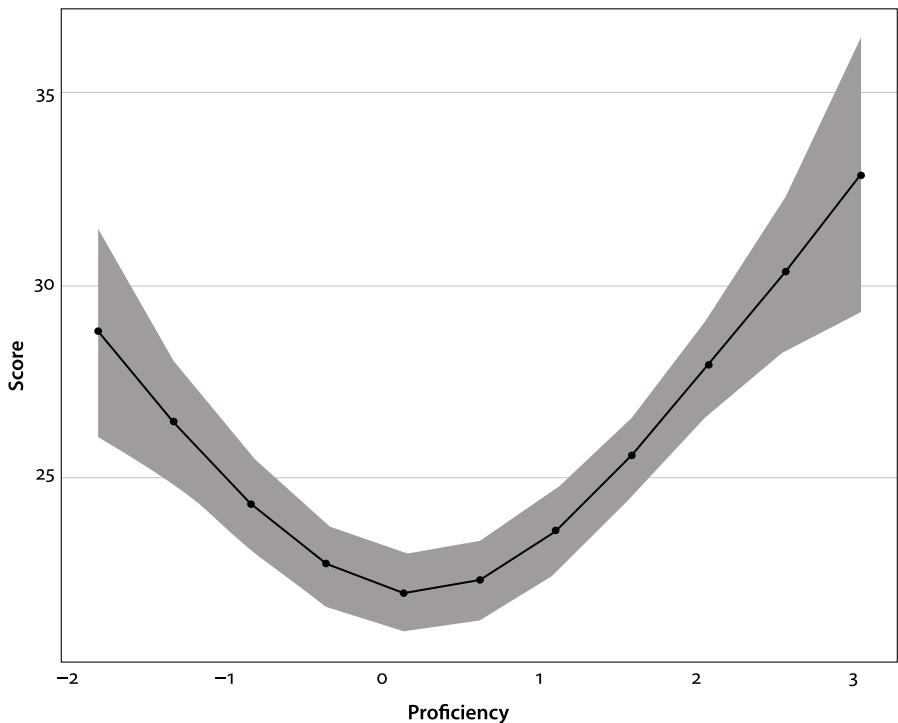
Before looking at the summary of the `gam_1` model, we plot the predicted effect of proficiency. For this, we will use the `tidygam` package (Coretta, 2023), which provides users with utility functions that make extracting and plotting predictions from GAMs more accessible.

```
# Load tidygam package
library(tidygam)

# Extract model predictions from 'gam_1', assign them to 'gam_1_preds'
gam_1_preds <- predict_gam(gam_1)

# Plot predictions
plot(gam_1_preds, series = "proficiency")
```

Figure 2 is a plot of the predicted effect of proficiency on accuracy, based on the `gam_1` model. We will look into the details of `predict_gam()` shortly, but for now, notice that accuracy initially decreases as proficiency increases. At a proficiency score of approximately 0.5, accuracy begins to increase, resulting in the typical U-shaped curve described above. At this juncture we will inspect the model summary.



**Figure 2.** Predicted accuracy score depending on proficiency, from a GAM (gam\_1 model)

```
# Print a summary of the model 'gam_1'
summary(gam_1)

Family: gaussian
Link function: identity

Formula:
score ~ s(proficiency)

Parametric coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 24.559     0.331   74.19 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:
        edf Ref.df    F p-value
s(proficiency) 2.922 3.675 16.46 <2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) =  0.242  Deviance explained = 25.3%
GCV = 22.352  Scale est. = 21.914  n = 200
```

The relevant parts of the summary for the time being are the Parametric coefficients: table and the Approximate significance of smooth terms: table. The former contains the estimate of the intercept. This intercept represents the same intercept you would estimate in a standard linear model. In our case, the estimate of the intercept is the predicted accuracy score when proficiency is equal to zero. Thus our interpretation according to the summary, is that when proficiency is 0, learner accuracy is about 25.

Importantly, what we are interested in is the effect of proficiency on accuracy, rather than the estimate of accuracy when proficiency is zero, i.e., the intercept. Information on the effect of proficiency is found in the second table, which contains estimates of the smooth terms. Alas, these estimates are not informative about the effect, per se, but rather indicate if the relationship between proficiency and accuracy is linear (or not) in nature. More specifically, the `edf` estimate, which stands for Estimated Degrees of Freedom, is equal to one in the case of a perfectly linear relationship. A value greater than one is an indication of non-linear effects. Crucially, the EDF estimate is not informative with regard to the exact shape of the effect. The only way to assess this is to plot the model predictions, as we have done in Figure 2.

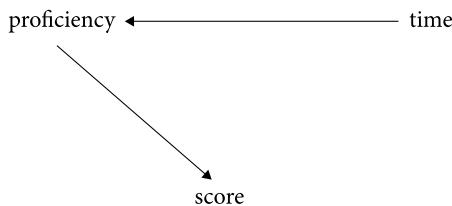
In our `gam_1` model, the `edf` of the smooth term for proficiency is 2.9, thus suggesting that the effect of proficiency on accuracy is not linear (Note: the `Ref.df`, reference degrees of freedom, and `F` values merely serve the purpose of being used to derive a *p*-value). The *p*-value indicates the probability of the observed smooth under the null hypothesis that the smooth is a horizontal flat line (in other words, under the null hypothesis that the variable of the smooth has no effect on the outcome variable, Simpson (2023)).

If we were to report this model and interpret the results, we could write something along the following lines: We fitted a Generalised Additive Model to learner accuracy scores, with a smooth term over proficiency (to model non-linear effects). According to the model, the effect of proficiency is non-linear and significant ( $F=16.46$ ,  $p < .0001$ ). Based on the prediction plot, we observe that at lower proficiency, learning scores initially decrease and then start increasing again when proficiency reaches approximately 0.5.

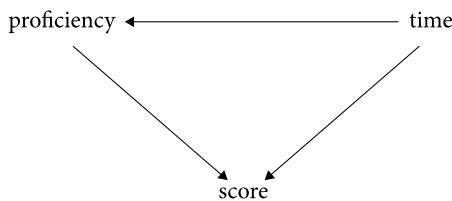
We began our analysis with a simplified model examining the effect of proficiency on accuracy at a single time point for pedagogical purposes. In the following section we will expand our analysis by refitting the data including time as a predictor.

### 3.4 Multiple smooth terms

As mentioned previously, the data `dat1` contains accuracy and proficiency scores from 200 subjects, taken at 10 time points. The data was simulated so that proficiency increased over time (at different degrees for different participants), as is often the case in language acquisition processes. An interesting question to consider is whether learning scores improve with time independently of proficiency, or if proficiency alone is causing learning scores to improve. We can approach this question using principles of causal inference. To this end, we develop directed acyclic graphs (DAGs) to describe the causal relationships involved. Though a full treatment of causal inference is beyond the scope of this paper, we believe the DAG is informative in aiding understanding of the model we will build. Readers interested in a deep dive into causal inference and DAGs are referred to McElreath (2019) and Pearl (2010).



(a) No direct effect of time on learning scores



(b) Direct effect of time on learning scores

**Figure 3.** Directed Acyclic Graphs for time, proficiency and learning score

The DAGs in Figure 3 represent the causal relationships between accuracy, proficiency scores and time in two scenarios. In (a), time affects proficiency and proficiency affects accuracy. In other words, time has a direct effect on proficiency, but not on accuracy; accuracy scores can be predicted from proficiency alone. In (b), on the other hand, time affects proficiency, as in (a), but also accuracy (and proficiency affects accuracy scores as in (a)). The practical interpretation is that time affects accuracy in two ways: through its effect on proficiency and directly through its effect on accuracy.

DAGs allow us to make causal statements based on statistical results. In this case, when including both time and proficiency as predictors in a GAM, time should not have an effect on learner accuracy if scenario (a) is correct (while it should have an effect if scenario (b) is correct). Now we fit the model.

```
# Fit model, assign to 'gam_2'
gam_2 <- gam(
  formula = score ~ s(proficiency) + s(time) + s(subj, bs = "re"),
  data = dat1
)
```

In `gam_2` we fit a GAM to learner accuracy `score` with two predictors: a smooth term over proficiency and a smooth term over time. We also include random effects to account for nested data from multiple participants. The syntax for random effects in GAMs is different from the more familiar syntax used in `lme4`. With `gam()`, we can specify random effects using smooth terms and the argument `re` (for Random Effects) basis function. Random intercepts are added with the syntax `s(INTERCEPT_EFFECT, bs = "re")` and random slopes with the syntax `s(INTERCEPT_EFFECT, SLOPE_EFFECT, bs = "re")`. Here we only include a *by*-subject random intercept for illustration, i.e., `s(subj, bs = "re")`. Below we inspect the model summary.

```
# Print summary of model 'gam_2'
summary(gam_2)

Family: gaussian
Link function: identity

Formula:
score ~ s(proficiency) + s(time) + s(subj, bs = "re")

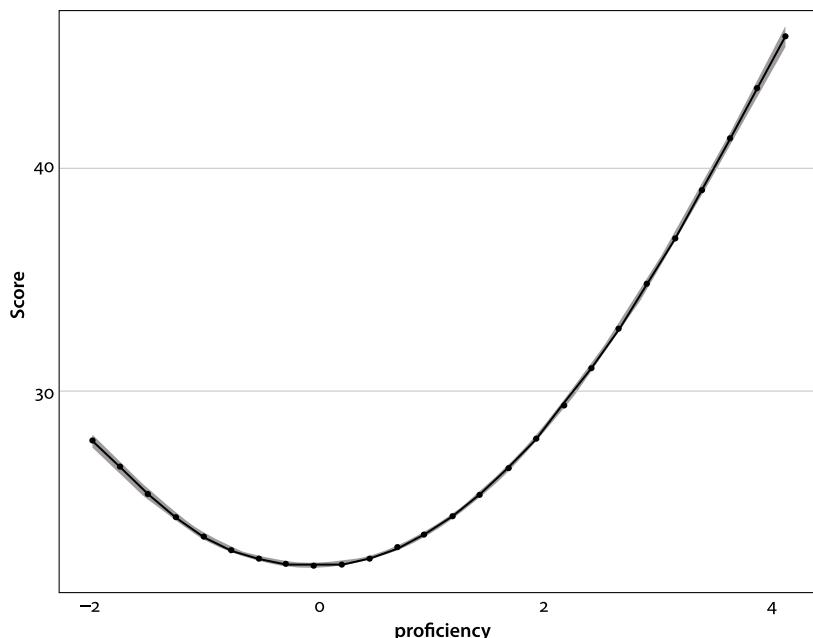
Parametric coefficients:
Estimate Std. Error t value Pr(>|t|)
(Intercept) 24.738     0.441   56.09 <2e-16 ***
---
Signif. codes: 0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 ' ' 1

Approximate significance of smooth terms:
          edf Ref.df      F p-value
s(proficiency) 8.699 8.976 4638.38 < 2e-16 ***
s(time)        1.000 1.000  20.27 7.36e-06 ***
s(subj)         198.943 200.000 2637.63 < 2e-16 ***
---
Signif. codes: 0 '****' 0.001 '***' 0.01 '**' 0.05 '*' 0.1 ' ' 1

R-sq.(adj) = 0.997 Deviance explained = 99.7%
GCV = 0.10023 Scale est. = 0.090681 n = 2200
```

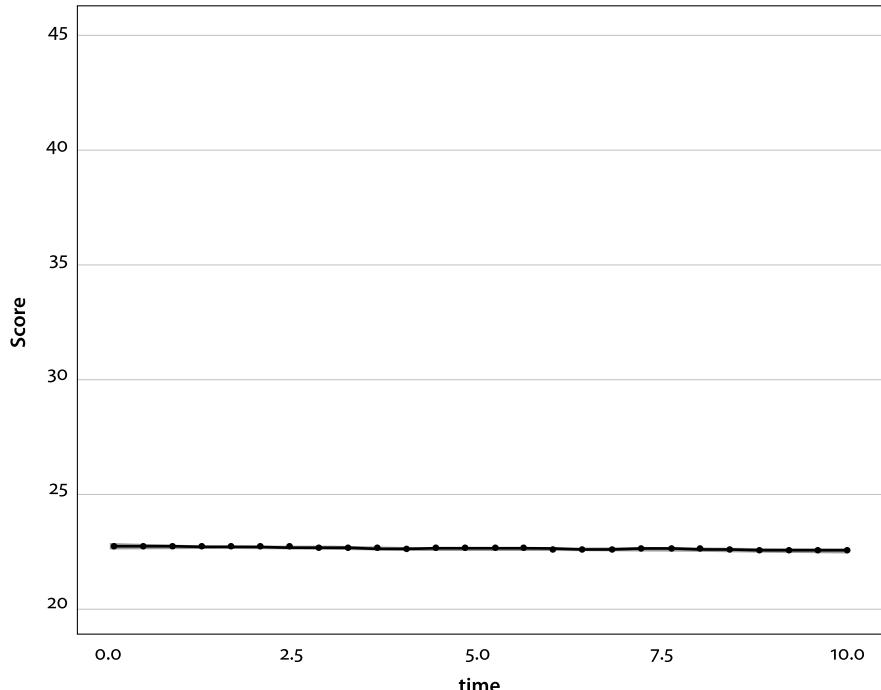
Now we plot the model predictions of accuracy score by proficiency and time point separately.

```
# Get model predictions for 'proficiency', assign them to 'gam_2_preds'  
gam_2_preds <- predict_gam(  
  gam_2, length_out = 25,  
  series = "proficiency",  
  exclude_terms = c("s(subj)"))  
)  
  
# Plot predictions  
plot(gam_2_preds)
```



**Figure 4.** Predicted accuracy score as a function of proficiency, from a GAM (gam\_2\_preds model)

```
# Get model predictions for 'time', assign them to 'gam_2_preds_t'  
gam_2_preds_t <- predict_gam(  
  gam_2, length_out = 25,  
  series = "time",  
  exclude_terms = c("s(subj)"))  
)  
  
# Plot predictions  
plot(gam_2_preds_t) +  
  ylim(20, 45)
```



**Figure 5.** Predicted accuracy score as a function of time, from a GAM (`gam_2_preds_t` model)

Figure 4 clearly shows a non-linear effect of proficiency on accuracy scores. On the other hand, Figure 5 indicates that, once accounting for the effect of proficiency, time has virtually no effect on accuracy, as expected based on the data simulation and DAGs. Note that according to the model the effect of time is linear but *non-zero*, hence the significant  $p$ -value for that smooth. The effect is, nonetheless, practically zero for two reasons: (1) first, the data was simulated in such a way that time has no effect; (2) second, when inspecting the predicted effect one can see that the mean difference in accuracy score between time point 1 and 11 is about 0.2, which on the simulated accuracy scale is negligible. This constitutes a nice caveat about issues with significance testing (see Gigerenzer (2004) for more on the topic). Generally, we advise researchers to rely more on model plotting to assess effects.

We could report this model like so: We fitted a Generalised Additive Model to learner accuracy scores, with a smooth term over proficiency (to model non-linear effects of proficiency) and a smooth term over time (to model non-linear effects of time). We also included a random intercept smooth for subject (to account for overall difference in mean accuracy by subject). According to the model, the effect of proficiency is non-linear and significant ( $F = 4638.38$ ,  $p <$

.0001). Based on the prediction plot, we observe that at lower proficiency, learning scores initially decrease and then start increasing again when proficiency reaches approximately 0.5. While the model summary indicates a significant effect of time on accuracy, the effect is negligibly small, which can be observed in the prediction plot of time. To summarise, the model suggests a non-linear effect of proficiency on learners' accuracy scores, but no effect of time.

In this example, we examined u-shaped development in second language acquisition. Specifically, we explored a simulated data set and fit a series of models, increasing in complexity. Our modelling strategy was outlined *a priori* using a DAG. We illustrated how to incorporate non-linear effects into the model via smoothing terms using the function `s()` inside the call to `gam()`. Next, we focused on using the fitted object to generate and plot predictions using `predict.gam()` and, finally, we walked through how one can report the resulting analysis. In the next section, we will consider vowel production in bilingual speech.

#### 4. Case study 2: L2 vowel production in simultaneous and late bilinguals

To motivate our next example, we will now consider the production of Spanish vowels in three groups of bilinguals: simultaneous/native English/Spanish bilinguals (henceforth “es”) and late learners of Spanish with low and high levels of proficiency, `es_0` and `es_1`, respectively. Spanish has five phonemic vowels, /i, e, a, o, u/, and is often described as being a syllable-timed language in which the spectral envelope is rather stable and stressed and unstressed syllables have approximately the same duration. Though there is evidence that unstressed vowels can centralize to a certain degree, the differences between stressed and unstressed vowels are believed to be imperceptible (Martínez Celdrán, 1984). American English, on the other hand, has a larger vowel inventory and is often described as a stress-timed language in which the duration between stressed syllables is approximately the same. Vowel reduction can occur in unstressed syllables, which typically manifests via shortening and/or centralization, with unstressed vowels often reducing to schwa ([ə]). Accordingly, English speaking learners of Spanish face a substantial obstacle when it comes to producing and perceiving Spanish vowels (See Aldrich, 2014; Bland, 2016; Cobb, 2009; Cobb & Simonet, 2015; Iruela, 1997, among others). To wit, they often display cross-linguistic influence by producing unstressed Spanish vowels with [ə], particularly in the case of /a/, e.g., “casa” (Eng. *house*) as [ka.sə], and diphthongizing vowels in final position (Cobb & Simonet, 2015).

In the following example, we consider the production of Spanish vowels in the aforementioned groups of bilingual participants. Interestingly, to our knowledge, all of the current research investigating the cross-linguistic influence of vowel reduction processes in learners of Spanish have utilized modelling strategies that rely upon formant measurements at a single time point, typically the mid-point of the vowel, rather than scrutinizing the entire trajectory of the spectral envelope (see de la Fuente Iglesias & Pérez Castillejo (2022) and Amengual & Simonet (2020)). In our view, the dynamic nature of vowel formants poses an interesting use case for modelling using non-linear methods, such as GAMs. To this end, we have simulated phonetic data from three aforementioned groups of bilinguals: simultaneous English/Spanish bilinguals, beginner late English/Spanish bilinguals and advanced late English/Spanish bilinguals. The data contains the Euclidean Distance (EuD) of three Spanish vowels /a, i, u/ from the vowel space centroid, taken from word-final non-stressed vowels, along nine equidistant time points. The measure is a good proxy for vowel reduction phenomena, which is to be expected in late bilinguals.

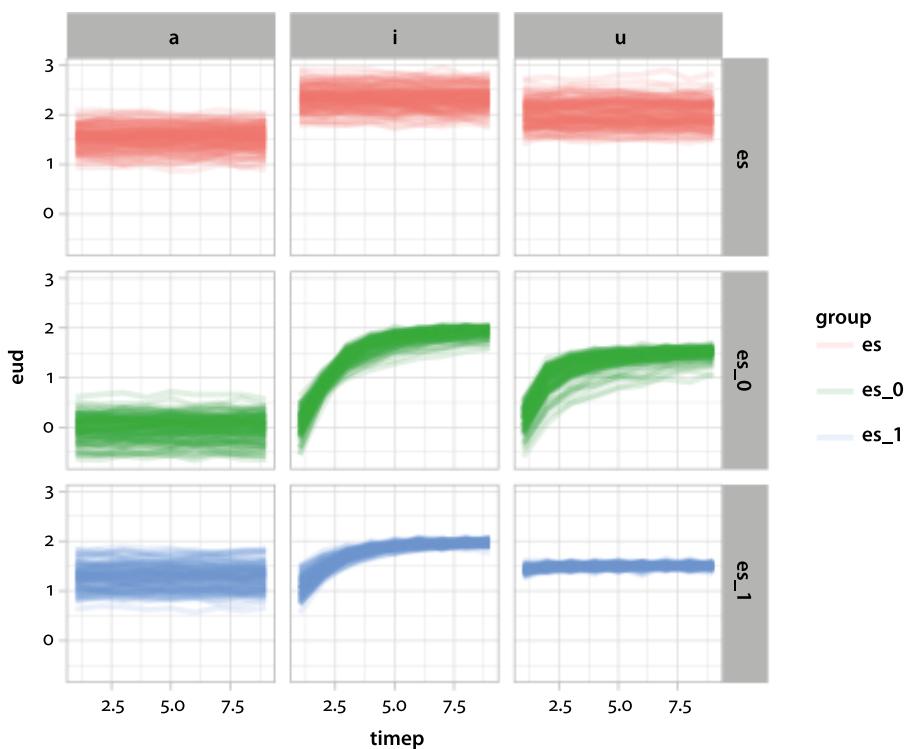
```
# Load data set
dat2 <- readRDS("data/dat2.rds")

# Print first 6 rows
head(dat2)

# A tibble: 6 × 7
  subj group vowel rep   eud timep id
  <fct> <chr> <chr> <int> <dbl> <int> <chr>
1 es_1 es     i      1  1.85    1 es.1.1
2 es_1 es     i      1  1.86    2 es.1.1
3 es_1 es     i      1  1.83    3 es.1.1
4 es_1 es     i      1  1.87    4 es.1.1
5 es_1 es     i      1  1.89    5 es.1.1
6 es_1 es     i      1  1.81    6 es.1.1
```

The above code prints the first six rows of the data frame. Specifically, we observe a subset of the synthetic data of a simultaneous English/Spanish bilinguals (`es_1`). It can be observed in Figure 6 that for the `es` group (top row), the Euclidean distance is quite stable across the duration of the vowel for all three vowels and away from `o` (which corresponds to a mid-central vowel), indicating no reduction takes place (as is to be expected for Spanish). On the other hand, in the second and third row of Figure 6, we can see some vowel reduction at play: for /a/, beginner late bilinguals (`es_o`, middle row) completely reduce the vowel to a mid-central vowel, while the advanced late bilinguals (`es_1`, bottom row) reduce to a lesser degree; /i/ and /u/ are produced with a diphthongal quality, where

the first part of the vowel is quite reduced. In sum, comparing beginners and advanced bilinguals shows that the latter have less reduction, presumably a by-product of increasing proficiency in Spanish.



**Figure 6.** Euclidean Distance through 9 time points taken from within vowels

We can use a GAM to model the Euclidean distance along time points in the different vowels and different bilingual groups. Fitting smooth terms for different levels of categorical predictors can be achieved by specifying the categorical predictor as the value of the `by` argument inside the smooth term function `s()`. For example, `s(timep, by = vowel)` would ensure a non-linear effect of time point is estimated for each vowel. In our data, we want to model the interaction of vowel (/a, i, u/) and group (es, es\_0, es\_1). Unfortunately, the additive nature of GAMs does not allow the direct specification of interactions between smooth terms like one would do in a linear model (interactions require product operations). Note that interactions between parametric terms are supported, but these allow only for interactions in the overall height of the smoothers, not in their shape. The method presented in this section allows users to specify “interactions” in smooth terms that can model differences in both height and shape.

To include interactions between smooth terms, we can simply construct a new predictor with the combination of vowel and group and use that as the `by`-variable in the smooth terms. Due to how smooth terms are constructed when you include a `by`-variable, it is necessary to also include the variable as a parametric term (i.e., a classical linear term). We would also like to allow for the modelled trajectories to vary by subject both in height and shape. While random effects included with `bs = "re"` (which we used in the first case study) account for difference in height and rotation of the fitted curves, *random factor smooth interaction* terms fit a different curve to each level of the chosen factor (in our case, subject). Random factor smooth interaction terms are specified with `bs = "fs"`: the first argument in `s()` should be the variable to smooth over (here, `timep`) and the second argument is the factor (here `subj`), the levels of which will be fitted separately. Finally, the `m` argument corresponds to the order of the smoothing penalty: higher orders lead to more smoothing (the default is 2, a second-order penalty). Setting `m = 1` corresponds to a higher constraint on the smooths of each level in the random factor, and it is thus a safeguard against over-fitting.

```
# Generate group_vowel column to model the interaction
dat2 <- dat2 |>
  mutate(
    vowel = as.factor(vowel),
    group_vowel = interaction(group, vowel),
  )

# Fit model
gam_3 <- gam(
  formula = eud ~
    # Parametric term for group_vowel
    group_vowel +
    # Smooth term with group_vowel by-variable
    s(timep, by = group_vowel, k = 9) +
    # Random factor smooths
    s(timep, subj, bs = "fs", m = 1),
  data = dat2
)

# Print model summary
summary(gam_3)

Family: gaussian
Link function: identity

Formula:
eud ~ group_vowel + s(timep, by = group_vowel, k = 9) + s(timep,
  subj, bs = "fs", m = 1)
```

```

Parametric coefficients:
Estimate Std. Error t value Pr(>|t|)    
(Intercept) 1.533926  0.023247 65.983 < 2e-16 ***
group_voweles_0.a -1.539712  0.032877 -46.833 < 2e-16 ***
group_voweles_1.a -0.255481  0.032877 -7.771 8.59e-15 ***
group_voweles.i 0.791406  0.007385 107.157 < 2e-16 ***
group_voweles_0.i -0.052418  0.032877 -1.594 0.111  
group_voweles_1.i 0.222520  0.032877  6.768 1.38e-11 ***
group_voweles.u 0.485860  0.007385 65.786 < 2e-16 ***
group_voweles_0.u -0.318526  0.032877 -9.689 < 2e-16 ***
group_voweles_1.u -0.044394  0.032877 -1.350 0.177  
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Approximate significance of smooth terms:
          edf Ref.df      F p-value    
s(timeep):group_voweles.a 1.000 1.000 0.112 0.73770  
s(timeep):group_voweles_0.a 1.000 1.000 0.117 0.73226  
s(timeep):group_voweles_1.a 1.000 1.000 0.023 0.87934  
s(timeep):group_voweles.i 1.000 1.000 0.021 0.88372  
s(timeep):group_voweles_0.i 6.221 7.201 1770.941 < 2e-16 ***
s(timeep):group_voweles_1.i 5.245 6.298 468.819 < 2e-16 *** 
s(timeep):group_voweles.u 1.000 1.000 0.009 0.92287  
s(timeep):group_voweles_0.u 7.568 7.945 777.603 < 2e-16 *** 
s(timeep):group_voweles_1.u 2.484 3.086 4.621 0.00269 **  
s(timeep,subj)           56.008 534.000 6.025 < 2e-16 *** 
---
Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

R-sq.(adj) = 0.939 Deviance explained = 94%
GCV = 0.029735 Scale est. = 0.029455 n = 97201

```

Let's extract the model predictions and inspect them.

```

# Get model predictions, assign them to object 'gam_3_preds'
gam_3_preds <- predict_gam(gam_3, exclude_terms = "s(timeep,subj)")
gam_3_preds

# A tibble: 99 × 6
  group_vowel timeep     eud      se lower_ci upper_ci
  <fct>       <dbl>    <dbl>    <dbl>    <dbl>    <dbl>
1 es.i        1     2.32   0.0155   2.29     2.35
2 es.u        1     2.02   0.0155   1.99     2.05
3 es.a        1     1.53   0.00809  1.52     1.55
4 es_0.i      1     0.0930  0.0473   0.000345 0.186
5 es_0.u      1     0.154   0.0476   0.0611   0.248
6 es_0.a      1    -0.00855  0.0418  -0.0888   0.0717
7 es_1.i      1     1.10   0.0470   1.00     1.19
8 es_1.u      1     1.45   0.0447   1.36     1.53
9 es_1.a      1     1.28   0.0410   1.20     1.36
10 es.i       1.8    2.32   0.0139   2.30     2.35
# 89 more rows

```

The function `predict_gam()` automatically samples 10 equidistant points from numeric variables (like `timep`, based on the default `length_out = 10` argument) and all levels in categorical variables (like `group_vowel`). For each combination of the sampled points and levels, the function returns the value of the outcome (`eud`) and the standard error (`se`). The lower (`lower_ci`) and upper (`upper_ci`) boundaries of the 95% Confidence Interval are also returned.

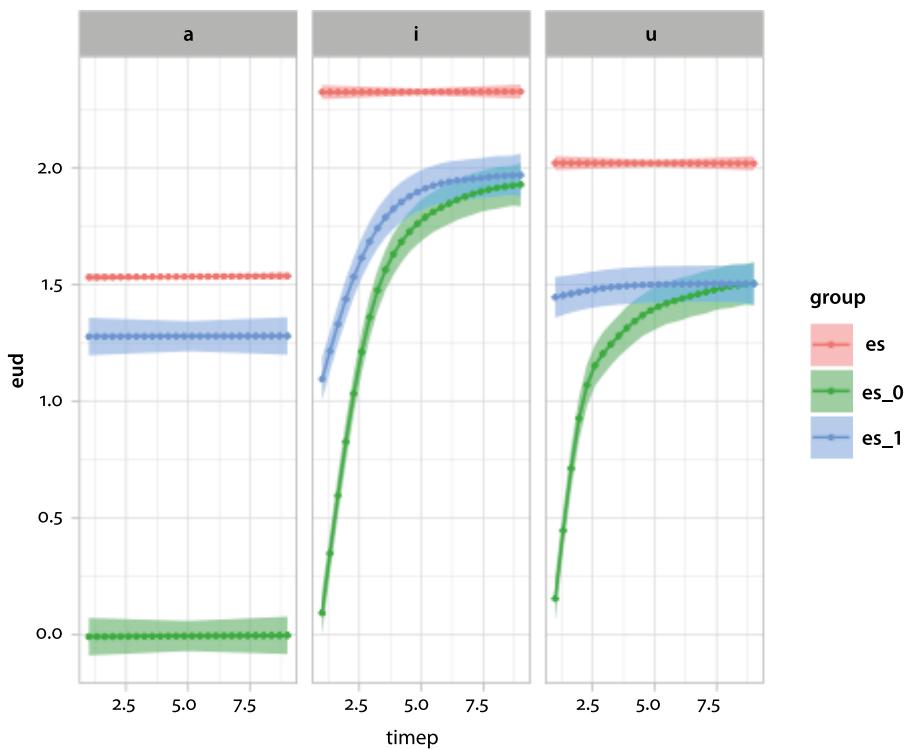
Furthermore, `predict_gam()` returns an object of class `tidygam` which can be plotted with `plot()` (similarly to how `ggpredict()` from the `ggeffects` package works, Lüdecke (2018)). However, we might want to do some processing of the predictions before plotting: in particular, we might want to split the `group_vowel` column into the two original variables, `group` and `vowel`. We can achieve this straight from the '`predict_gam()`' function. Additionally, we can also extract more than 10 time points to get a smoother curve.

```
# Get model predictions, separate group_vowel column
gam_3_preds_2 <- predict_gam(
  gam_3,
  length_out = 25,
  exclude_terms = "s(timep,subj)",
  separate = list(group_vowel = c("group", "vowel")))
)
gam_3_preds_2

# A tibble: 234 × 7
  group vowel timep     eud      se lower_ci upper_ci
  <chr> <chr> <dbl>    <dbl>   <dbl>    <dbl>   <dbl>
1 es     i     1     2.32  0.0155   2.29    2.35
2 es     u     1     2.02  0.0155   1.99    2.05
3 es     a     1     1.53  0.00809  1.52    1.55
4 es_θ   i     1     0.0930 0.0473  0.000345 0.186
5 es_θ   u     1     0.154  0.0476  0.0611  0.248
6 es_θ   a     1    -0.00855 0.0410  -0.0888  0.0717
7 es_1   i     1     1.10  0.0478  1.00    1.19
8 es_1   u     1     1.45  0.0447  1.36    1.53
9 es_1   a     1     1.28  0.0410  1.20    1.36
10 es    i    1.32  2.32  0.0148  2.30    2.35
# [i] 224 more rows
```

The syntax `list(group_vowel = c("group", "vowel"))` instructs the function to split the `group_vowel` column into two columns, `group` and `vowel`. Splitting is done based on the `sep_by` argument, which is a period (.) by default. Now we can plot the predictions. (Figure 7):

```
# Plot predictions
gam_3_preds_2 |>
  plot(series = "timep", comparison = "group") +
  facet_grid(cols = vars(vowel))
```



**Figure 7.** Predicted Euclidean distance along the duration of the vowel, for /a, i, u/ in simultaneous bilinguals (es), beginners late bilinguals (es\_0) and advanced late bilinguals (es\_1)

What follows is an example of how the model and results could be reported in prose.

We fitted a Generalised Additive Model to Euclidean Distance values taken along the duration of the vowels /a, i, u/ in simultaneous, low-proficiency late and high-proficiency late English/Spanish bilinguals. We included a parametric term for the variable group\_vowel which is the combination of group (es, es\_0, es\_1) and vowel (a, i, u) to model mean differences in EUD between group/vowel, and a smooth term over time point (timep) with group\_vowel as a by-variable to model non-linear trajectories of EUD in each group/vowel condition. We have also included a random factor smooth by subject over time point to model by-subject differences in mean EUD and trajectory shapes.

The simultaneous English/Spanish bilingual data (group es) in Figure 7 clearly shows stable EUD throughout the duration of vowels /a/, /i/ and /u/, indicating that no unstressed vowel reduction occurs (of course, since the data is simulated,

the EUD trajectories are unnaturally flat). Moreover, the EUDs for each vowel differ thus indicating that each vowel is distinguished despite absence of stress as we would expect in Spanish. When looking at the beginner late bilinguals (group es\_0), the plot indicates that /a/ is reduced to [ə] (the EUD is 0), that /i/ and /u/ are strongly diphthongised, and that the first part of the diphthong is close to [ə] while the second part does not quite reach the EUD values of the simultaneous English/Spanish bilinguals, again as we would expect in English speakers learning Spanish.

In the high-proficiency late English/Spanish bilinguals we can observe almost no reduction in /a/ (the EUD values are close to the simultaneous bilinguals' values and flat) and a lesser degree of reduction in /i/ and /u/, although the EUD value for the second part of the diphthongised vowels still don't quite reach the value of the simultaneous bilinguals. Moreover, the second part of /i/ and /u/ is very similar in both late bilingual groups, thus suggesting that while there is less diphthongisation in the high-proficiency group, this group still hasn't developed simultaneous bilingual-like targets.

## 5. Summary

This introductory tutorial covered the basics of fitting non-linear effects using Generalised Additive Models (GAMs) and the mixed-effects counterpart Generalised Additive Mixed Models (GAMMs) with the mgcv and tidygam packages in R.

Two literature-based simulated case studies were employed to illustrate how to fit and plot GAMs. Case study 1 was about U-shaped learning and focused on modelling non-linear effects with smooths (`s()` terms in R syntax) and on obtaining model predictions for plotting. Case study 2 illustrated a time-series analysis of euclidean distance trajectories of three vowels (/a, i, u/) of three bilingual groups (simultaneous English/Spanish, low-proficiency late, high proficiency late). This case study introduced readers to *by*-variables, used to model effects of categorical factors on smooths, including interactions between factors, and random factor smooth interactions, which fit random effects.

This tutorial on its own will not be sufficient to develop an adequate GAM analysis but constitutes a low barrier-to-entry for readers who are interested in learning about GAMs. We encourage those who wish to run a full GAM analysis to peruse the tutorials mentioned in the introduction: Sóskuthy (2017), Sóskuthy (2021), Pedersen et al. (2019), Tamminga et al. (2016), and Wieling (2018). Important topics that were not covered here are selection of type and number of basis functions in smooths, autocorrelation and autoregressive models, statistical inference and modelling other distribution families.

## 6. Data availability statement

The compendium with code and data used in this tutorial can be found at [https://github.com/stefanocoretta/gamm\\_biling](https://github.com/stefanocoretta/gamm_biling).

## Funding

Open Access publication of this article was funded through a Transformative Agreement with University of Edinburgh.

## References

- Aldrich, A. C. (2014). Acquisition of L2 phonology: An acoustic analysis of the centralization of L2 Spanish /a/ in adult L1 English-speaking learners. MA Thesis. Brigham Young University.
-  Amengual, M., & Simonet, M. (2020). Language dominance does not always predict cross-linguistic interactions in bilingual speech production. *Linguistic Approaches to Bilingualism*, 10(6), 847–872.
- Bland, J. (2016). Speech style, syllable stress, and the second-language acquisition of Spanish /e/ and /o/. MA thesis, Virginia Polytechnic Institute; State University.
-  Botezatu, M. R., Misra, M., & Kroll, J. F. (2024). Proficiency in a second language influences processing of print-to-sound mappings. *Linguistic Approaches to Bilingualism*, 14(3), 285–309.
-  Carlucci, L., & Case, J. (2013). On the necessity of U-shaped learning. *Topics in Cognitive Science*, 5(1), 56–88.
- Casani, E. (2020). Valutare la competenza morfosintattica in italiano L2. Una validazione corpus-based dei livelli del QCER. In E. Nuzzo, E. Santoro, & I. Vedder (Eds.), *Valutazione e misurazione delle produzioni orali e scritte in italiano lingua seconda* (pp. 15–26). Franco Cesati Editore.
-  Cazden, C. B. (1968). The acquisition of noun and verb inflections. *Child Development*, 39(2), 433–448.
- Cobb, K. (2009). La pronunciación de vocales átonas en español: La aplicación de reglas fonológicas por parte de hablantes no-nativos del español. MA Thesis. The University of Arizona.
-  Cobb, K., & Simonet, M. (2015). Adult second language learning of Spanish vowels. *Hispania*, 98(1), 47–60.
- Coretta, S. (2023). *Tidygam: Tidy prediction and plotting of Generalised Additive Models*. [\url{https://CRAN.R-project.org/package=tidygam}](https://CRAN.R-project.org/package=tidygam)
-  de la Fuente Iglesias, M., & Pérez Castillejo, S. (2022). L1 phonetic permeability and phonetic path towards a potential merger: The case of Galician mid vowels in bilingual production. *Linguistic Approaches to Bilingualism*, 12(2), 191–219.

- Ervin, S. M., & Miller, W. R. (1963). Language development. In H. W. Stevenson, J. C. Kagan, C. C. Spiker, N. B. Henry, & H. G. Richey (Eds.), *Child psychology: The sixty-second yearbook of the National Society for the Study of Education, Part 1* (pp. 108–143). University of Chicago Press.
- doi** Gass, S., & Selinker, L. (2008). *Second language acquisition. An introductory course*. Routledge.
- Geeslin, K. L., & Guijarro-Fuentes, P. (2006). A longitudinal study of copula choice: Following development in variable structures. In N. Sagarra & A. J. Toribio (Eds.), *Selected procedures of the 9th hispanic linguistics symposium* (pp. 144–156). Cascadilla Proceedings Project.
- doi** Gigerenzer, G. (2004). Mindless statistics. *The Journal of Socio-Economics*, 33(5), 587–606.
- doi** Hastie, T., & Tibshirani, R. (1986). Generalized Additive Models. *Statistical Science*, 1(3), 297–310.
- Iruela, A. (1997). Adquisición del vocalismo español por holandeses: Análisis en estilo semiespontáneo. *Estudios de Fonética Experimental IX*, 135–180.
- Kellerman, E. (1985). If at first you do succeed... In S. M. Gass & C. G. Madden (Eds.), *Input in Second Language Acquisition* (pp. 345–353). Newbury.
- Lightbown, P. (1983). Exploring relationships between developmental and instructional sequences in L2 acquisition. In H. Seliger & M. Long (Eds.), *Classroom oriented research in second language acquisition* (pp. 217–243). Newbury House.
- doi** Long, M. H. (1990). The least a second language acquisition theory needs to explain. *TESOL Quarterly*, 24(4), 649–666.
- doi** Lüdecke, D. (2018). Ggeffects: Tidy data frames of marginal effects from regression models. *Journal of Open Source Software*, 3(26), 772.
- doi** Marcus, G. F., Pinker, S., Ullman, M., Hollander, M., Rosen, T. J., Xu, F., & Clahsen, H. (1992). Overregularization in language acquisition. *Monographs of the Society for Research in Child Development*, 57(4), i–178.
- Martínez Celadrán, E. (1984). *Fonética*. Teide.
- McElreath, R. (2019). *Statistical rethinking: A Bayesian course with examples in R and Stan*. Boca Raton, FL: CRC Press.
- doi** Pearl, J. (2010). An introduction to causal inference. *The International Journal of Biostatistics*, 6(2).
- doi** Pedersen, E. J., Miller, D. L., Simpson, G. L., & Ross, N. (2019). Hierarchical generalized additive models in ecology: An introduction with mgcv. *PeerJ*, 7, e6876.
- doi** Pliatsikas, C., & Marinis, T. (2013). Processing empty categories in a second language: When naturalistic exposure fills the (intermediate) gap. *Bilingualism: Language and Cognition*, 16(1), 167–182.
- R Core Team. (2023). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. <https://www.R-project.org/>
- doi** Rumelhart, D. E., & McClelland, J. L. (1986). On learning the past tenses of English verbs. Implicit rules or parallel distributed processing? In J. L. McClelland & D. E. Rumelhart (Eds.), *Parallel distributed processing: Explorations in the microstructure of cognition*. MIT Press.

- Shirai, Y. (1990). U-shaped behavior in L2 acquisition. *Variability in Second Language Acquisition: Proceedings of the Tenth Meeting of the Second Language Research Forum*, 2, 685–700.
- Simpson, G. (2023). *What do p-values mean in generalized additive model?* Cross Validated.
- Sóskuthy, M. (2017). *Generalised additive mixed models for dynamic analysis in linguistics: A practical introduction*. arXiv.org preprint, arXiv:1703.05339.
-  Sóskuthy, M. (2021). Evaluating generalised additive mixed modelling strategies for dynamic speech analysis. *Journal of Phonetics*, 84, 101017.
-  Tamminga, M., Ahern, C., & Ecay, A. (2016). Generalized additive mixed models for intraspeaker variation. *Linguistics Vanguard*, 2(s1), 33–41.
-  Wieling, M. (2018). Analyzing dynamic phonetic data using generalized additive mixed modeling: A tutorial focusing on articulatory differences between L1 and L2 speakers of English. *Journal of Phonetics*, 70, 86–116.
-  Williams, S., Guijarro-Fuentes, P., & Vulchanova, M. (2022). U-shaped trajectories in an L2 context: Evidence from the acquisition of verb morphology. *Vigo International Journal of Applied Linguistics*, 19, 223–266.
- Winter, B. (2013). *Linear models and linear mixed effects models in R with linguistic applications*. arXiv:1308.5499. <https://arxiv.org/pdf/1308.5499.pdf>
-  Winter, B. (2020). *Statistics for linguists: An introduction using R*. Routledge.
-  Wood, S. (2006). *Generalized additive models: An introduction with R*. CRC Press.
-  Wood, S. (2011). Fast stable restricted maximum likelihood and marginal likelihood estimation of semiparametric generalized linear models. *Journal of the Royal Statistical Society (B)*, 73(1), 3–36.

## Address for correspondence

Stefano Coretta  
 Linguistics and English Language  
 The University of Edinburgh  
 3 Charles St  
 Edinburgh EH8 9AD  
 United Kingdom  
 s.coretta@ed.ac.uk  
 <https://orcid.org/0000-0001-9627-5532>

## Co-author information

Joseph V. Casillas  
 Department of Spanish and Portuguese  
 Rutgers University  
 joseph.casillas@rutgers.edu  
 <https://orcid.org/0000-0001-8735-9910>

## Publication history

Date received: 20 December 2023

Date accepted: 21 October 2024

Published online: 29 November 2024